

Is the Mars Curiosity Rover taking a selfie?

Scott Armstrong
DSI 2020, 9 June 2020


Dataset

Deep Mars: CNN Classification of Mars Imagery for the PDS Imaging Atlas

**Kiri L. Wagstaff,¹ You Lu,¹ Alice Stanboli,¹
Kevin Grimes,¹ Thamme Gowda,^{1,2} Jordan Padams¹**

¹Jet Propulsion Laboratory, California Institute of Technology, Pasadena, CA 91109, {firstname.lastname}@jpl.nasa.gov

²Information Sciences Institute, University of Southern California, Marina Del Rey, CA 90292, tg@isi.edu




Caltech



Images




Processing




imageTrain.csv

```
id, pixel 1, pixel 2, ...  
0, 124, 103, ...  
1, 114, 148, ...  
...
```



Distributed Training

Scala (Java) → Spark is cross-platform



Master
Runs main process, distributing requests among workers

Worker
Completes requests of master and sends back result

Results

MLP NN → 11% accuracy

Linear SVM → 17% accuracy

Random Forest → **75%** accuracy

- 256 trees

- max depth 16 nodes

Next Steps

Improve images

- Color
- Larger images
- Augmentation

Improve distribution

- Distributed file management
- Image processing in Scala



Thank you